An *ab-initio* Diagrammatic Monte Carlo prototype for the LiF electron polaron

Samuel Recker

August 10, 2026

Abstract

This note fixes the **goals**, deliberate **non-goals**, and **deferred** extensions of the `rmc-ab-initio` project: a diagrammatic Monte Carlo (DiagMC) prototype, built on the `ironmc` engine, that replaces the analytic Fröhlich electron–phonon vertex $|g(\mathbf{q})|^2 \propto \alpha/q^2$ with the *real* electron–phonon matrix elements $g(\mathbf{k}, \mathbf{q}, \nu)$ obtained from density-functional perturbation theory and Wannier interpolation (Quantum ESPRESSO + EPW). The first validation material is the electron polaron in LiF, a wide-gap ionic insulator that sits close to the textbook Fröhlich limit and therefore provides a built-in analytic cross-check. The scope below is the *first* prototype; Section 4 records what is intentionally postponed.

1 Overview

`ironmc` is a Rust DiagMC engine that samples the imaginary-time self-energy expansion of a polaron. Its only physics application today is the *analytic* Fröhlich model, whose vertex and dispersions are closed-form scalars. This project adds a second, *ab-initio* model: the same sampler, but with a material-specific electron–phonon (e–ph) vertex read from a real DFT/DFPT + Wannier calculation.

The data pipeline is

$$\underbrace{\text{pw.x, ph.x}}_{\text{Quantum ESPRESSO}} \longrightarrow \underbrace{\text{epw.x}}_{\text{Wannier e–ph tensors}} \longrightarrow \underbrace{\text{Rust reader + reconstruction}}_{\texttt{crates/apps/rmc-ab-initio}} \longrightarrow \text{DiagMC.}$$

DFPT supplies the phonons, the Born effective charges $\mathbf{Z}^*_\kappa$, and the electronic dielectric tensor $\boldsymbol{\varepsilon}_\infty$; EPW supplies the e–ph coupling in the localized Wannier representation $g(\mathbf{R}_e, \mathbf{R}_p)$, from which $g(\mathbf{k}, \mathbf{q}, \nu)$ is recovered by an inverse Fourier transform.

Why LiF. Rocksalt LiF has a single, non-degenerate conduction-band minimum at Γ, one dominant longitudinal-optical (LO) branch, and an intermediate-to-strong Fröhlich coupling $\alpha \approx 5$. It is therefore the cleanest first material: nearly the textbook Fröhlich case, with a genuine analytic limit to validate against.

Units. We work in Fröhlich units $\hbar = m^* = \omega_{\mathrm{LO}} = 1$. In these units the parabolic electron dispersion is exactly $p^2/2$ and the Einstein LO phonon is exactly 1, so the *only* material-specific quantity in the parabolic prototype is the vertex $|g(\mathbf{q})|^2$. The Fröhlich coupling is

$$\alpha \;=\; \frac{e^2}{\hbar}\sqrt{\frac{m^*}{2\hbar\omega_{\mathrm{LO}}}}\left(\frac{1}{\varepsilon_\infty} - \frac{1}{\varepsilon_0}\right), \tag{1}$$

and energies are reported back in meV via $\hbar\omega_{\mathrm{LO}}$.

2 Goals

G1 — Real vertex. Drive the DiagMC sampler with the *ab-initio* e–ph vertex in place of the analytic one. The vertex is decomposed as

$$g_{mn\nu}(\mathbf{k},\mathbf{q}) \;=\; g^{\mathcal{S}}_{mn\nu}(\mathbf{k},\mathbf{q}) \;+\; g^{\mathcal{L}}_{mn\nu}(\mathbf{k},\mathbf{q}), \tag{2}$$

the short-range (Wannier-interpolated) part plus the long-range polar part.

G2 — Engine untouched. Reuse the state-generic `ironmc` engine (`Runner`, `Update`, `MetropolisKernel`) *bit-for-bit*. In Fröhlich units the mass and phonon-energy constants never change, so the ab-initio model enters through a single vertex-weight seam rather than an engine rewrite.

G3 — Correct polar long range. Restore the long-range Fröhlich $1/q$ divergence that dominates the LiF polaron. EPW was run with `lpolar=.true.`, which *subtracts* the long-range part before Wannierization, so $g^{\mathcal{S}}$ from the tensor is short-ranged and $g^{\mathcal{L}}$ must be added back analytically (Section 1, Eq. (3)).

G4 — Clean Rust reader. Implement the EPW-file readers and the Wannier reconstruction from scratch in Rust, with no Python at runtime. The existing Python post-processor (`patt126-ab-initio-diagmc`) is used as a *reference only*, not copied; the rewrite also fixes a latent index assumption in the reference.

G5 — Two-front validation. Benchmark the LiF electron polaron against *both* (i) Fröhlich-limit DiagMC energetics at α_{LiF} (Mishchenko *et al.*, 2000), and (ii) *ab-initio* LiF polaron numbers (Sio, Verdi, Poncé, Giustino, 2019).

The long-range term used for G3 (Verdi–Giustino form, at the Γ slice) is

$$g^{\mathcal{L}}_{\nu}(\mathbf{q}) \;=\; i\,\frac{4\pi e^2}{\Omega}\sum_{\kappa}\sum_{\mathbf{G}\neq-\mathbf{q}}\left(\frac{\hbar}{2M_\kappa\omega_{\nu\mathbf{q}}}\right)^{1/2}\frac{(\mathbf{q}+\mathbf{G})\cdot\mathbf{Z}^*_\kappa\cdot\mathbf{e}_{\kappa\nu}(\mathbf{q})}{(\mathbf{q}+\mathbf{G})\cdot\boldsymbol{\varepsilon}_\infty\cdot(\mathbf{q}+\mathbf{G})}\,e^{-i(\mathbf{q}+\mathbf{G})\cdot\boldsymbol{\tau}_\kappa}\,e^{-\frac{(\mathbf{q}+\mathbf{G})\cdot\boldsymbol{\varepsilon}_\infty\cdot(\mathbf{q}+\mathbf{G})}{4\Lambda}}, \tag{3}$$

with Ω the cell volume, $M_\kappa, \boldsymbol{\tau}_\kappa$ the atomic masses and positions, $\mathbf{e}_{\kappa\nu}$ the phonon eigenvectors, and Λ an Ewald range-separation parameter. As $q \to 0$ this behaves as $1/q$, reproducing the Fröhlich tail $|g_F(\mathbf{q})|^2 \propto \alpha/q^2$.

3 Non-goals (deliberate scope of the first prototype)

These are conscious simplifications, chosen so the prototype is small, exact where it matters, and cheaply validated. Each is lifted in Section 4.

N1 — Electron polaron only. The single, non-degenerate conduction-band minimum at Γ. The hole/valence polaron (degenerate top valence band) is out of scope.

N2 — Γ-slice vertex. We approximate $g(\mathbf{k},\mathbf{q}) \approx g(\Gamma,\mathbf{q})$, i.e. the vertex depends on the phonon momentum $\mathbf{q}$ only, evaluated as if the electron sits at the band bottom. True electron-momentum dependence is a non-goal here.

N3 — One Einstein LO branch. Keep only the Fröhlich-active LO branch and treat it as dispersionless at ω_{LO}. The remaining branches and the dispersion $\omega_{\mathrm{LO}}(\mathbf{q})$ are not modeled.

N4 — **No tensor compression.** The reduced object the sampler needs is the scalar field $|g(\Gamma, \mathbf{q}, \mathrm{LO})|^2$, which is small on any grid. It is precomputed once and interpolated; *no* SVD/Tucker low-rank compression is used. (Fixing $\mathbf{k} = \Gamma$ is itself an exact, lossless dimensional reduction that supersedes lossy compression in this regime.)

N5 — **Reweighting, not native sampling.** The ab-initio vertex is introduced by reweighting the existing analytic Fröhlich importance-sampling envelope, not by a native $|g_{\mathrm{ab}}(\mathbf{q})|^2$ momentum sampler.

N6 — **No engine parameterization.** The engine's mass/phonon-energy constants stay fixed at unity (Fröhlich units); no `MASS`/`OMEGA` genericization is introduced for the parabolic prototype.

N7 — **No GPU path.** The ab-initio model targets the CPU driver only; the GPU (`#[cube]`) mirror is not extended to it.

N8 — **Single material.** LiF is the sole validation target; this is not yet a general multi-material framework.

4 Deferred to future work

Listed roughly in the order they would be relaxed. None of these require reworking the Phase-A readers or reconstruction; they extend the same pipeline.

D1 — **Full $g(\mathbf{k}, \mathbf{q})$ k-dependence.** Relaxing the Γ-slice (N2) turns the vertex into a function on the six-dimensional $(\mathbf{k}, \mathbf{q})$ space, which cannot be precomputed onto a grid. This is exactly the regime where **SVD/Tucker compression + on-the-fly evaluation** of the Wannier tensor becomes necessary (reversing N4). Relevant at strong coupling, where the electron explores momenta away from Γ.

D2 — **Multi-branch, dispersive phonons.** Include all phonon branches and the momentum dependence $\omega_\nu(\mathbf{q})$, replacing the single Einstein LO branch (N3).

D3 — **Non-parabolic electrons.** A non-parabolic $\varepsilon(\mathbf{k})$ requires genericizing the engine's dispersion (the deferred `MASS`/`OMEGA` parameterization and `Model`-trait seam), lifting N6.

D4 — **Native q-sampler.** A momentum sampler drawn directly from $|g_{\mathrm{ab}}(\mathbf{q})|^2$, replacing the reweighting of N5 for sampling efficiency.

D5 — **Hole polaron.** The degenerate valence-band-maximum polaron (lifting N1), which needs a multi-band, degenerate treatment.

D6 — **GPU path.** A `#[cube]` GPU implementation of the ab-initio model (lifting N7), guarded by the existing lockstep parity tests.

D7 — **More materials.** Generalize beyond LiF once the LiF prototype is validated (lifting N8).

5 Validation ladder

The goals of Section 2 are checked in three rungs of increasing physical content:

1. **Refactor parity.** The Fröhlich model, re-expressed through the new ab-initio code path, reproduces the existing engine's $\alpha = 1$ benchmark ($E \approx -1.016$, $Z \approx 0.59$).

2. **Fröhlich-limit + unit bridge.** Fed the *analytic* $|g(\mathbf{q})|^2 \propto \alpha_{\mathrm{LiF}}/q^2$, the ab-initio model must match the pure Fröhlich engine at α_{LiF} against Mishchenko *et al.* (2000); and the *real* tensor's $|g(\Gamma, \mathbf{q}\to 0, \mathrm{LO})|^2$ must recover the same $1/q^2$ asymptote.

3. **Material comparison.** The real LiF vertex yields a polaron formation energy and effective mass compared to Sio *et al.* (2019).

6 Practical pitfalls and fixes (A0/A1)

Turning the EPW run into a *usable* vertex meant working through several failures whose symptoms were misleading. They are recorded here so they are not rediscovered; the blow-by-blow diagnosis, with reproducing commands, lives in `datasets/LiF/EPH-NAN-DIAGNOSIS.md`.

P1 — Ultrasoft pseudopotentials give an all-`NaN` e–ph tensor (root cause). The first A0 deck used ultrasoft pseudopotentials (SSSP-efficiency GBRV). EPW's electron–phonon matrix elements are only supported with *norm-conserving* pseudopotentials — the ultrasoft augmentation charge is not carried into the coupling. The coupling came out finite *only* at $\mathbf{q} = \Gamma$ (the trivial reference q, which needs no symmetry rotation) and `NaN` at every other $\mathbf{q}$; the Bloch→Wannier sum then spread the `NaN` to 100% of `LiF.epmatwp`. Because `ph.x` *does* support ultrasoft, the phonons, dynamical matrix, $\mathbf{Z}^*_\kappa$ and $\boldsymbol{\varepsilon}_\infty$ were all correct — only the e–ph coupling was wrong, which is why the fault was easy to misattribute. *Fix:* switch Li and F to norm-conserving Pseudo Dojo ONCV (scalar-relativistic PBE). This raises the required plane-wave cutoff (`ecutwfc` : 50 → 96 Ry, `ecutrho` : 400 → 384 Ry, i.e. the norm-conserving 4×), which makes `ph.x` ∼3–4× slower (the wavefunctions move from the small smooth FFT grid onto the dense grid) — an unavoidable cost of the correct pseudopotential family. With norm-conserving pseudos the tensor is finite at every $\mathbf{q}$, and $\boldsymbol{\varepsilon}_\infty$, $\mathbf{Z}^*_\kappa$ are essentially unchanged.

P2 — "Right size, wrong values": a silent failure. The broken run finished cleanly and wrote `LiF.epmatwp` at exactly the expected byte size; the only hint was an easily-dismissed `IEEE_INVALID` note at exit. The file was 100% uninitialised-memory `NaN` (a uniform garbage word that even *changed* between reruns). An earlier "A0 done" had been declared on file size alone, without ever checking the tensor *values*. *Fix:* the pipeline (`run_pipeline.nu`) now runs an explicit finiteness check on `LiF.epmatwp` and aborts loudly on `NaN`, so a broken run can no longer masquerade as a good one.

P3 — Two dead ends from assuming an upstream bug. Before the pseudopotential cause was found, two "EPW bug" hypotheses were pursued and *both failed*: (i) widening the Wannier outer window `dis_win_max` to enclose every band made the coarse tensor *worse*, not better; (ii) a bounds-guard patch to EPW's `gmap_sym` routine (a genuine latent out-of-bounds read) compiled in cleanly yet left the output byte-identical — so that read was not the cause. The lesson, recorded explicitly: suspect the *inputs* (here, an unsupported pseudopotential family) before patching upstream code.

P4 — The reference `.epb` reader has the wrong record layout. The Python reference (`patt126`) reads the coarse Bloch `.epb` file assuming a single band count throughout the record. The real record interleaves a *20-band* eigenvalue array with a *15-band* coupling array (of which only a 5×5 block is populated); a naïve parse silently mis-aligns and yields a *false positive* — a "the coarse data looks fine" reading that briefly sent the debugging down a wrong path. *Fix:* the true record layout was reverse-engineered and byte-verified, and the Rust reader uses it directly.

P5 — epw coarse-grid prerequisites (reproducibility). Several setup constraints are load-bearing and non-obvious: the `nscf` k-points must be an *explicit* Γ-centred crystal list (an automatic, half-shifted Monkhorst–Pack mesh breaks EPW's $\mathbf{k}+\mathbf{q}$ folding); the whole `scf/ph/nscf/epw` chain must use *one* Quantum ESPRESSO version (a mixed-version `phsave` corrupts the heap); and EPW needs a large stack (`ulimit -s unlimited`). Finally, changing the pseudopotentials shifts the band edges ($\sim$0.15 eV here), so the disentanglement windows (`dis_froz_max`, `dis_win_max`) must be re-tuned against the new `nscf` eigenvalues — otherwise the frozen conduction-band minimum, which must stay exact, can slip out of its window.